

Word Count: 7,971

Figures: 1

Tables: 3

Masculinity, Meaning, and Misogyny: Linguistic Markers of Intra- and Intergroup Identity Threat Responses in r/JordanPeterson

Eli Cochrane* & Mark J. Landau

University of Kansas, Department of Psychology, Lawrence, KS, USA

*** Correspondence:**

Eli Cochrane

elimcochrane@ku.edu

Keywords: political psychology, social identity, digital masculinities, manosphere, online radicalization, natural language processing, identity threat

Abstract

In recent decades, progressive social movements have disrupted masculine roles that traditionally afforded men a clear social identity, provoking a “crisis of masculinity” shown to fuel ideological extremism and sociopolitical polarization. Within this context, Jordan Peterson has emerged as an influential public figure esteemed by millions of men as a philosophical guide providing routes back to traditional masculinity. Despite his initial focus on self-improvement, Peterson’s online discourse community has been implicated in radical discourse that denigrates scapegoated outgroups, particularly women and LGBTQ+ individuals. The current research analyzes when and how Peterson-related discourse shifted as a case study in the radicalization of masculinity-focused discourse communities. We apply the source model of identity threat to distinguish between intergroup and intragroup threats and compare their differential impacts on outgroup prejudice. Using computational natural language processing techniques, we longitudinally analyzed the r/JordanPeterson subreddit ($N = 87,035$ posts; 2017-2024), the centralized forum for Peterson-related discourse. Using latent profile analysis of linguistic features (via LIWC-22 and tidyLPA) and transformer-based topic modeling (via BERTopic), we identified the community’s discursive profiles and thematic trends. Interrupted time series analysis tracked linguistic shifts surrounding major identity threats. Following an intragroup threat, the community exhibited gradual boundary enforcement toward outgroups, self-directed reactive hostility, and stagnant group cohesion. Conversely, an intergroup threat prompted less explicit outgroup discussion in favor of principle-based ideological defense and enhanced group cohesion. These findings extend the source model to a naturalistic digital environment, demonstrating that threat origin fundamentally shapes online radicalization trajectories. Ultimately, effective mitigation of digital extremism requires understanding these threat-specific dynamics and opportunities for intervention.

Introduction

Young men are experiencing a profound crisis of masculinity. For generations, men could organize their identity formation around culturally sanctioned norms and values associated with their expected societal roles of breadwinner, protector, and patriarch. Yet in recent decades, social movements for gender equality, racial justice, and LGBTQ+ rights have challenged traditional masculine norms, values, and roles. As such, many contemporary men feel that their path toward a culturally sanctioned masculine identity has been destabilized (Kimmel, 2017), motivating them to find or create new models of masculinity around which they can orient their identity formation.

In the digital age, this identity (re)construction increasingly takes shape in online male-centered discourse communities, collectively labeled the “Manosphere,” that provide space for like-minded men to discuss identity-relevant topics. Although sometimes entirely user-led, Manosphere communities are often spearheaded by charismatic “manfluencers” who seek to guide other men in their path to feeling stable in their masculinity (Renström & Bäck, 2024). Despite their seemingly innocuous intentions, manfluencers and their followers have been accused of fomenting misogyny and transphobia and propagating extreme right-wing social policies (Gottzén et al., 2026; Institute for Strategic Dialogue, 2022). The current research aims to shed light on this radicalization pattern by integrating theoretical insights and methodological advances that have thus far occupied separate corners of group-level psychology research.

At a theoretical level, we build on evidence that although ingroup identification is sufficient to produce some degree of outgroup stereotyping and prejudice (Tajfel & Turner, 2000), it is specifically when group identity is threatened that prejudice intensifies into more hostile forms (Branscombe et al., 1999). Understanding how masculinity communities respond to different types of threats may help explain when and why these communities sometimes radicalize to embrace outgroup hostility. To this end, we utilize Greenaway and Cruwys’s (2019) *source model* of identity threat, which distinguishes *intragroup* threats originating from within the group (e.g., internal dissent or norm violation) and *intergroup* threats originating from outgroups. As detailed shortly, prior research shows that intergroup threats increase prejudice, whereas the effects of intragroup threat on prejudice are mixed. Extending this work, we examine whether masculinity-based communities differentially denigrate women and LGBTQ+ individuals in response to intra- and intergroup threats.

At a methodological level, we model ongoing identity construction processes via natural language processing techniques that identify group members’ topical foci, discursive styles, and their relation over time (Tausczik & Pennebaker, 2010). These methods allow us to identify when and why masculinity discourse communities shift their focus from relatively depoliticized self-help topics to hostile outgroup stereotyping and vilification.

We focus on the online masculinity community inspired by Jordan B. Peterson, a prominent figure in the Manosphere. This community’s discourse features a mix of relatively depoliticized self-help content and outgroup hostility, particularly targeting women and LGBTQ+ individuals, providing a productive case study for the naturalistic analysis of extremism in response to intra- and intergroup threats.

Background

r/JordanPeterson as a Digital (Neo)Masculinity Community

Psychologist and self-help author Jordan B. Peterson is one of the most popular online manfluencers, having amassed nearly nine million subscribers and over a billion total views on YouTube. Unlike most other manfluencers, Peterson’s brand of neomascularity emerged not from

explicitly anti-feminist positions but from an ideological framework focused on self-improvement. His 2018 bestselling book, *12 Rules for Life*, propelled him to the cultural mainstream by providing men with academically framed practical advice for living well (Champagne, 2020).

On Reddit, a social media platform defined by its use of topic-based forums called “subreddits,” the subreddit r/JordanPeterson serves as the most active centralized forum for Peterson-related discourse. As of 2026, it had over 300,000 subscribers.

Around the 2018 release of *12 Rules for Life*, a cursory glance revealed that r/JordanPeterson posts largely focused on relatively depoliticized topics, including self-improvement strategies and philosophical approaches to life’s challenges. Peterson described himself as explicitly opposed to misogyny, instead wanting to focus on helping men without necessarily derogating women and other feminized outgroups.

In the years that followed, however, r/JordanPeterson has gained notoriety for promulgating radical transphobic, homophobic, and misogynistic content (Lafrance et al., 2025). Although many contributors continued to look to the forum for personal guidance (Nesbitt-Larking, 2022), many others embraced derogation of feminized others, namely, women and LGBTQ+ individuals which, at its most extreme, may devolve into violent behavior (Habib et al., 2022; Nicholas et al., 2026).

Peterson himself has shifted ideologically, as his *Conservative Manifesto* (2022) frames the masculinity crisis in explicitly right-wing terms. A full account of his fanbase’s ideological trajectory would model Peterson’s own radicalization, but we limit our scope to the potential links between r/JordanPeterson’s radicalization and salient events that may pose distinct identity threats. This event-driven approach explores the possibility that radicalization is not inevitable or ideologically predetermined, but may instead be precipitated by identifiable, and potentially intervenable, moments of perceived threat. Next, we review a theoretical framework for classifying threats and anticipating potential defensive responses.

Source-Relevant Identity Threat and Outgroup Derogation

Social identity theory (SIT; Tajfel & Turner, 2000) posits that individuals derive much of their self-esteem from their group memberships. Hence, threats to the ingroup’s status or distinctiveness undermine members’ self-esteem, motivating compensatory affirmation of the ingroup’s value. A common strategy for restoring a positive social identity is outgroup derogation, wherein group members devalue outgroups to elevate the ingroup’s relative status. In spaces like the Manosphere, where members may already feel a loss of social status, SIT suggests we would observe especially close identification with the ingroup and high prejudice in response to events that threaten the group’s collective self-esteem.

SIT has inspired other theoretical perspectives to zoom in on distinct identity-relevant threats and their compensatory strategies. Uncertainty-identity theory (UIT; Hogg, 2007) frames identity threat as a loss of certainty about one’s self-concept, an aversive state that people seek to minimize by embracing group identities that prescribe clear roles, norms, and other constructs for clarifying one’s self-concept. Similarly, optimal distinctiveness theory (ODT; Brewer, 1991) posits that group members’ desire to view the ingroup as distinct from other groups drives them to derogate outgroups as a means of affirming the ingroup’s uniqueness.

While SIT, UIT, and ODT distinguish between the *types* of identity threat, other lines of research emphasize the importance of the *source* of a threat in predicting group members’ compensatory strategies (Branscombe et al., 1999). In this vein, the source model of identity threat (Greenaway & Cruwys, 2019) distinguishes between *intergroup* and *intragroup* threats.

Intergroup threats originate externally, typically perceived as an outgroup-originated challenge to the ingroup's power, values, or sustained existence. Extensive research consistently demonstrates that these external threats, particularly symbolic threats to a group's worldview or moral framework, trigger a defensive "rally around the flag" effect that manifests in outgroup derogation and ingroup solidarity (Giles & Evans, 1985; Riek et al., 2006; Rios et al., 2018). To neutralize the psychological discomfort of the threat (whether based on self-esteem, certainty, positive distinctiveness, etc.), the ingroup unifies to reinforce its boundaries and attack the seemingly responsible outgroup.

One event in Peterson's public trajectory stands out as an intergroup threat. In June 2022, Peterson had posted on Twitter (now *X*) about transgender male actor Elliot Page, stating that Page had "just had her breasts removed by a criminal physician." After Twitter moderators flagged his post for violating policies against harassment and hateful conduct, Peterson was banned from the site and faced widespread public backlash. Despite this public outrage, Peterson defended his actions, stating on an Instagram post in collaboration with conservative news outlet *The Daily Wire* that he would "rather die" than delete his post (Peterson & *The Daily Wire*, 2022). Peterson framed his ban from Twitter as a coordinated, malicious assault by "woke moralists" against free speech and "biological reality." This event is a quintessential intergroup threat, since a group's central leader and its core ideological values were under direct, punitive attack by an external, institutional outgroup (i.e., Twitter moderation and the broader progressive ideology it represented). We note that real-world events are rarely pure instances of a single threat source, as a leader's public punishment may also seed internal doubt, so we treat each event as predominantly, rather than exclusively, one threat type.

The source model identifies intragroup threats as arising from within the group itself, disrupting internal cohesion and shared norms (Doosje et al., 1995). Such threats challenge group members' confidence that the group is a unified whole, frequently resulting in ideological schisms or loss of collective identity.

Peterson's fanbase experienced an intragroup threat around November 2019, when Peterson's public appearances became less common and he experienced a noticeable decline in health. These symptoms went unexplained until February 2020, when Peterson was revealed to have been struggling with substance abuse and seeking treatment internationally. The sudden and initially unexplained loss of leadership between November 2019 and February 2020 represents an exemplary intragroup threat, as his fanbase had to navigate their group identity independent of their philosophical anchor. This threat was likely heightened as the circumstances of Peterson's absence were due to mental health challenges not unlike those for which his followers sought his advice. Navigating this threat thus had the potential to fracture the group's identity, necessitating collective identity defense.

The literature remains somewhat divided on how intragroup threats influence defensive outgroup derogation. Because intragroup threats compromise the integrity of the ingroup, group members may turn their hostility inward to police boundaries within the group and punish deviant members (the so-called black sheep effect; Marques et al., 1988) or distance themselves from the group to protect their individual self-esteem (Ellemers et al., 1997).

Other findings suggest that intragroup identity threat increases outgroup derogation as a means of restoring group cohesion. Extensions of UIT predict that when the group fractures, the resulting uncertainty can motivate group members to scapegoat a familiar outgroup (Hogg, 2014; Hogg & Adelman, 2013; Sherman et al., 2007). Furthermore, in the absence of clear leadership, members may engage in competitive social signaling. By performing extreme, visible acts of

outgroup derogation, individuals can use radical rhetoric as an impression management strategy to prove their loyalty (Klein et al., 2007) and group prototypicality (Schmid & Muldoon, 2015).

Because empirical findings regarding the divergent responses to intergroup and intragroup threats are mixed, a crucial step is to observe whether and how these identity threat sources differentially predict outgroup derogation. The current research pursues this goal through a naturalistic analysis of the linguistic contours of r/JordanPeterson as a case study in the expanding, increasingly politicized cultural discourse surrounding masculinity.

Method

On text-based forums like Reddit, users are known to carefully construct their language to demonstrate communicative competence, mutual understanding, and social inclusion (Panek, 2022). As a result, researchers can fruitfully apply discourse analysis techniques to analyze the language used on such forums for insights into the salient topics and attitudes by which users collectively interpret their group, its value, and what it means to be a group member (Gee, 2005). We applied this approach to longitudinally analyze r/JordanPeterson using a theory-driven, inductive computational pipeline that uses natural language as a marker of individual and collective identities. This approach allows us to examine how contributors to this forum engage with various concepts relevant to their identity construction goals and how those discursive patterns evolve in response to events likely to be interpreted as identity threats.

This pipeline captures two complementary layers of the community's discourse, which we analyze both separately and in combination. The first layer concerns *how* members express themselves. To this end, we applied the 2022 version of the Linguistic Inquiry and Word Count software (LIWC-22; Boyd et al., 2022), a widely used dictionary-based psycholinguistic tool, to extract theoretically relevant linguistic features from each post. Latent Profile Analysis (LPA) then consolidated these features into a small set of discourse profiles, that is, recurring modes of expression that are relatively independent of the post's topical focus.

The second layer concerns *what* members discuss. Transformer-based topic modeling (via BERTopic) categorized the thematic content of the corpus into interpretable themes. Because any theme can, in principle, be discussed in any style, profiles and themes carry distinct information, and analyzing them separately reveals shifts in the community's style and its substance, while combining them reveals which linguistic styles represent each theme (e.g., whether discussions of women are marked by hostile reactivity or impersonal boundary enforcement). Finally, Interrupted Time Series (ITS) analysis tracked discursive changes in themes, discourse profiles, their combinations, and a metric of group affiliation following the identity-relevant events of current interest, alongside long-term trends across the subreddit's lifespan.

Data Collection

Data collection utilized archives of Reddit data acquired using the Pushshift.io API and provided for public access by Reddit user stuck_in_the_matrix (2025). These archive files contain complete historical records of subreddit activity and user interactions. From this online database, we collected all posts in r/JordanPeterson from every full year available at the time of collection ($N = 209,500$; Jan. 1, 2017 – Dec. 31, 2024). The resulting dataset included each post's title, text content, author username, date posted, comments, and an engagement score computed by subtracting "downvotes" (or "dislike" reactions) from upvotes ("like" reactions). As all posts are publicly available online, this research was exempt from Institutional Review Board review.

Labeling LIWC-Based Discourse Profiles

The corpus was first processed using LIWC-22, which operationalizes psychological constructs present in text data by counting occurrences of specific word types, capturing aspects

of language such as affective tone, psychological foci, and pronoun usage (Boyd et al., 2022; Tausczik & Pennebaker, 2010). This approach quantifies cognitive states through structural word choice rather than self-reported beliefs and has been validated as a robust predictor of diverse behavioral and psychological outcomes (Eichstaedt et al., 2018; Koutsoumpis et al., 2022; Pennebaker et al., 2003).

To minimize multicollinearity, we selected nine theoretically relevant LIWC variables. Specifically, we analyzed pronoun usage (first-person singular, first-person plural, and third person) and cognitive and affective processes most theoretically relevant to our research question (power, insight, certitude, anxiety, anger, and moralization).

To identify how users framed their posts, we conducted Latent Profile Analysis (LPA) using the tidyLPA package in *R* (Rosenberg et al., 2018). LPA is a statistical modeling technique that identifies unobserved, distinct subgroups (latent profiles) within a dataset based on patterns of continuous indicator variables. By applying LPA to LIWC variables, we can empirically group texts that share mathematical similarity based on their selected linguistic features. This profile-based approach, rather than analyzing LIWC variables in isolation, captures holistic discursive patterns; mapping typologies via LPA provides an ecologically valid representation of the group's communication styles, while also reducing high dimensional text into interpretable behavioral trends. Researchers have increasingly used this approach to cluster online users into profiles based on psycholinguistic features, using data from Reddit (Adelina et al., 2023; Yan et al., 2024), Twitter (Zarate et al., 2023), and Weibo (Han et al., 2022), with applications to diverse social-psychological states and behaviors.

Because the true number of linguistic subgroups in a corpus is unknown *a priori*, model selection in LPA requires an iterative approach. We estimated a series of models specifying between 1 and 10 distinct profiles to determine the optimal number of subgroups. The 5-profile model was selected based on a combination of optimal statistical fit indices: a minimized Bayesian Information Criterion (BIC), high classification certainty, and adequate minimum class sizes. To interpret and define each subgroup, we calculated standardized z-scores of each LIWC variable and qualitatively reviewed representative text samples. This process yielded five distinct discourse profiles, named for their core psycholinguistic signatures (Table 1). Briefly, Anxious Self-Focus denotes personal, worried self-disclosure; Hostile Reactivity denotes short, raw bursts of anger; Third-Person Focus denotes discussion of other people and groups at analytic distance; Group Baseline denotes the community's default collective register; and Impersonal Policing denotes directive, rule-based boundary enforcement.

Labeling Thematic Content

Whereas LIWC analyses were used to identify discourse profiles, we analyzed posts for topical themes using BERTopic (Grootendorst, 2022), a topic modeling pipeline implemented using the Python-based Sentence Transformers library (Reimers & Gurevych, 2019). Posts with fewer than ten words were excluded to ensure sufficient semantic content for accurate analysis. Sentence Transformers converted each post into vector embeddings that capture contextual semantic meaning, with dimensionality reduction using UMAP. Then, the HDBSCAN algorithm inductively grouped the posts into latent topics; to maximize topic assignment without compromising accuracy, we utilized BERTopic's internal probability calculations to reassign soft outliers back into existing clusters if their assignment probability exceeded a threshold of 0.10.

To synthesize these findings into human-interpretable themes, we qualitatively grouped the algorithm's confident topic clusters following reflexive thematic analysis methods (Braun & Clarke, 2006). This procedure collapsed the highly granular topics into ten overarching topics,

which were used to train a supervised classification pipeline to categorize remaining unassigned outlier posts.

Semantic embeddings were generated for the outlier texts using the all-MiniLM-L6-v2 transformer model (Wang et al., 2020) and passed through a trained Linear Support Vector Classifier (SVC) configured with balanced class weights. Because an SVC decision function calculates the geometric distance from the classifying plane, a value of 0.0 represents the exact mathematical boundary. Any predictions falling below this 0.0 threshold were classified as noise, indicating the text lacked sufficient semantic similarity to be positively assigned to any defined theme. Posts that met the word-count criterion but could not be confidently assigned to one of the ten substantive themes were retained under a residual “Other” category. Although unclassified discourse has no single topical interpretation, it indexes the community’s broader conversation beyond the defined themes; shifts within it therefore indicate whether a dynamic was confined to specific topics or characterized community discourse at large. The ten substantive themes (Table 2) span discussions of Peterson himself, psychology, philosophy, self-help registers (positive, negative, and improvement-focused), politics, and three outgroup categories (women, LGBTQ+ individuals, and political outgroups). After exclusions for word count and classification fit into a theme and linguistic profile, the final sample was $N = 87,035$.

Combining Profiles and Themes

Having defined the discourse profiles and discursive themes, we treated them as separable layers of the same posts. Profiles and themes were analyzed separately to track independent changes in the community’s expressive style and topical focus, and jointly to measure the psychological meaning of stylistic shifts (e.g., rising hostility within self-help discussions signals a different group process than rising hostility in discussions of outgroups). Crossing the two layers therefore measures how changes reflect ongoing identity construction processes as shown through natural language. As a descriptive baseline, we first examined which linguistic profiles were most often present in discussions of each theme, visualized in Figure 1.

Finally, we included a measure of Unquestioning Affiliation, a LIWC-based index developed by Ashokkumar and Pennebaker (2022). Scores were calculated by subtracting each post’s standardized cognitive processing score from its standardized group affiliation score, capturing demonstrated identification with the group. This metric is a powerful predictor of group cohesion, co-occurring with both self-reported identity fusion and behavioral measures of long-term group identification. Changes in Unquestioning Affiliation can therefore reflect average group identification demonstrated within each theme and linguistic profile.

Temporal Analyses

After labeling each post, we analyzed thematic and linguistic distributions in relation to four major points in Jordan Peterson’s public life, each likely to hold formative implications for the group’s collective identity: (1) the January 2018 release of *12 Rules for Life*, an initial positive identity construction event; (2) his medical withdrawal beginning in November 2019, here modeled as an intragroup threat; (3) his public reappearance in June 2020, taken to be an identity reconstruction event; and (4) his June 2022 Twitter ban, here modeled as an intergroup threat. Prior research informs expectations of how the group might respond discursively to each event type, outlined in Table 3.

To evaluate the impact of each event and test these predictions, we aggregated the corpus into continuous monthly intervals. For each outcome variable, we estimated an ITS model using generalized least squares (GLS) with a first-order autoregressive AR(1) error structure and maximum likelihood estimation to account for temporal autocorrelation in monthly observations.

Each model was fit on a 12-month window centered on the event of interest (six months prior, the month of the event, and the five months following) and included three predictors: a linear time counter capturing the pre-existing baseline trajectory, a dummy variable capturing any immediate level shift the month of the event, and a post-event time counter measuring the monthly rate of change following the event.

For each window, we analyzed changes in (a) relative theme volume, (b) relative linguistic profile volume, (c) theme-by-profile interaction proportions (the distribution of each profile's posts across themes), and (d) Unquestioning Affiliation. Outcome variables were expressed as monthly proportions. For theme and profile models, this was the share of all posts within a given month belonging to a given theme or linguistic profile, and for theme-by-profile models, it was the share of a given linguistic profile's monthly posts belonging to a given theme. Unquestioning Affiliation was operationalized as a continuous monthly index reflecting average scores, modeled across the community. Because this index is a standardized difference score rather than a proportion, its effects are reported in index units rather than percentage points. To characterize the community's baseline longitudinal trajectory independent of any specific event, we additionally estimated a macro-level trend model for each outcome variable using a single GLS regression with an AR(1) error structure fit across the full 96-month observation period.

This computational pipeline provides a multifaceted way to observe identity negotiation amid identity-relevant events and threats. By tracking thematic focus, linguistic framing, and Unquestioning Affiliation across the four events, we can empirically test the source model of identity threat and observe its capacity to explain group radicalization.

Results

The following sections explore discursive shifts within r/JordanPeterson across both the full observation period and the four identity-relevant events outlined in Table 3. We only report results that were significant and large to ensure clarity and conciseness. Overall longitudinal findings are included in reporting based on statistical significance alone ($p < .05$), as the full observation window provides sufficient power to detect genuine trends and the annualized shift is interpretable as a real-world rate of change.

Event-specific ITS findings are held to a stricter three-criterion inclusion threshold, specifically statistical significance ($p < .05$), a squared partial correlation (pr^2) $\geq .25$, equivalent to a partial correlation of .50 (a large effect by Cohen's [1988] benchmarks), and a minimum shift of 0.50 percentage points, because the short 12-month window and proportional outcome variables create conditions for inflated significance that the additional criteria can address. All shifts are reported as percentage point changes in a variable's monthly proportion relative to its pre-existing baseline trajectory. The exception is Unquestioning Affiliation, which is reported in standardized index units. Findings involving unclassified discourse are reported alongside the substantive themes and are interpreted as shifts in the community's broader conversation rather than in any particular topic. Interpretations were supported by our qualitative readings of random example posts within the corpus.

Overall Changes

Analysis of r/JordanPeterson's long-term trajectory reveals a marked ideological shift. Over the eight-year observation period, discourse moved away from the community's foundational topics and toward politicized outgroup derogation, providing baseline empirical support for the radicalization pathway within the Manosphere.

Drift away from founding identity. Discourse shifted away from topics that began as core to the group's identity, including Jordan Peterson himself (-0.64% per year) and psychology

(−0.56% per year), with Group Baseline framing following the same pattern in both (−0.77% and −0.49% per year, respectively). In their place, politics and LGBTQ+ outgroup discussions grew (+0.76% and +0.26% per year), and Group Baseline framing migrated accordingly (+0.71% and +0.34% per year) while declining in self-improvement (−0.33% per year). These findings suggest a progressive abandonment of the group’s founding intellectual and self-help topics in favor of politicized, outgroup-related content.

Increased hostility toward gendered outgroups. Hostile Reactivity toward women rose +1.97% per year, and Anxious Self-Focus increased toward both women and LGBTQ+ individuals (+0.19% and +0.09% per year, the latter a small shift in absolute terms). Impersonal Policing, language high in power and moralization and low in personal disclosure, declined toward women (−0.83% per year) and political outgroups (−1.05% per year) while rising within general political discourse (+2.26% per year) and Third-Person Focus (+0.60% per year). These findings suggest that outgroup hostility became more emotionally raw and less ideologically structured over time, with boundary-enforcing language increasingly used to shape the ingroup’s own political identity rather than to discipline outgroups directly.

Rising group cohesion. Unquestioning Affiliation rose by 0.059 index units per year, reflecting a sustained increase in affiliative language and a decrease in cognitive processing, consistent with UIT’s prediction that highly entitative groups provide clear norms that reduce identity ambiguity for self-uncertain individuals.

I. Community Formation: 12 Rules for Life

The release of *12 Rules for Life* triggered a massive surge of new members. This period represents the community’s formative stage, in which foundational discourse norms and collective identity were established. The data suggest that early identity formation was not built on outgroup hostility; rather, overt misogyny was suppressed while members policed identity boundaries through political and philosophical topics.

Idealization of Peterson’s ideas. Peterson-focused discourse, which had already been growing (+0.70% per month), spiked immediately (+3.30%) then declined rapidly (−1.46% per month), suggesting initial novelty gave way to belief internalization. These discussions were largely authoritative: Impersonal Policing toward Peterson spiked +15.23% immediately, while Group Baseline framing, which had been declining before the event (−2.85% per month), grew as a sustained trend (+4.93% per month). Third-Person Focus in discussions of Peterson, which had been growing (+1.15% per month), began to decline (−2.15% per month). These findings together suggest that although the fanbase initially idealized Peterson as a figure, they increasingly internalized his framework, engaging with him less as a topic in his own right. Unclassified discourse also dropped immediately (−2.70%), consistent with conversation concentrating within the community’s defined topics during this period.

Group cohesion. Unquestioning Affiliation grew steadily ($b = 0.024$ index units per month). Group Baseline framing around LGBTQ+ individuals, which had been declining before the event (−1.09% per month), spiked immediately (+4.72%), suggesting that discussions of outgroups were briefly folded into the community’s collective self-definition. These findings indicate that group cohesion formed quickly during this period, potentially centered on group definition in contrast to outgroups.

Low outgroup hostility. The book’s release was accompanied by a suppression of hostility toward women. Discussions of women rose immediately (+1.73%), yet Group Baseline framing toward women, which had been elevated before the event (+6.33% per month), declined as a sustained trend (−7.76% per month), as did Impersonal Policing toward women (−6.00% per

month), while Third-Person Focus on women spiked (+2.18%). These findings suggest that although the community discussed women more, these discussions were marked by psychological distance rather than norm-enforcing or collective identity construction. The period of mainstream popularity may have created an initial identity centered on Peterson's life advice and philosophy that was strong enough that hostility toward feminized outgroups was not needed to build group identity.

Political and philosophical boundary creation. Political and philosophical antagonism became immediately common. Hostile Reactivity toward political outgroups grew as a sustained trend (+2.28% per month). In philosophical discussions, Hostile Reactivity reversed from declining before the event (−1.64% per month) to spiking immediately (+3.64%) and continuing as a trend (+1.06% per month), while Group Baseline framing of philosophical topics collapsed (−20.70%) even as philosophy grew as a topic (+0.70% per month). These findings indicate that the community engaged with Peterson's ideas reactively from the outset, frequently deploying them to define themselves against ideological others.

Failure-based anxiety and externalization. Overall Anxious Self-Focus dropped immediately (−4.31%). Across self-help topics, Anxious Self-Focus declined in improvement-focused discussions (−4.03% per month) while growing in failure-focused (+2.24% per month) and philosophical discussions (+0.87% per month). Self-improvement declined as a topic (−2.00% per month) while failure-focused content grew (+0.91% per month). Third-Person Focus shifted in parallel, declining in improvement-focused content (−1.40% per month) while growing in failure-focused content (immediately +1.42%, then +0.55% per month) and psychology (+0.71% per month). These findings indicate that when discussing failure, members increasingly referenced other people rather than themselves, suggesting Peterson's psychological frameworks were deployed outward to attribute blame rather than inward for self-improvement.

II. Intragroup Threat: Peterson's Withdrawal

Peterson's abrupt disappearance in late 2019 removed the community's epistemic anchor under circumstances that may have been viewed as contradicting his core message of self-mastery. Analysis of this event tests the source model's prediction that intragroup threats undermine group cohesion, resulting in compensatory outgroup derogation. The data show that outgroup hostility did emerge, but through gradual, impersonal norm enforcement rather than reactive discharge, and without converting into group cohesion.

Inward focus following loss of identity anchor. Peterson-related discussions, which had been growing before the event (+1.14% per month), dropped immediately (−6.11%). Hostile Reactivity and Third-Person Focus in Peterson-related discussions followed the same pattern, growing before the event (+2.89% and +1.37% per month, respectively) but collapsing immediately after (−19.51% and −6.94%), with Third-Person Focus continuing to decline (−1.17% per month). Anxious Self-Focus toward Peterson grew slowly (+0.62% per month). These findings suggest that Peterson's disappearance discouraged both analytic and hostile engagement with him, with the loss instead processed gradually over time. Meanwhile, unclassified discourse grew as a sustained trend (+1.36% per month), indicating that conversation increasingly diffused beyond the community's defined topics.

Compensatory political norm enforcement. Psychology rose as a topic (+2.12%) with increased third-person referencing (+2.23%), while Group Baseline framing in psychology, which had been growing (+2.15% per month), collapsed immediately (−8.10%). Impersonal Policing in politics reversed from declining (−1.35% per month) to spiking (+8.73%) as Anxious Self-Focus in politics dropped (−2.15%), and Impersonal Policing in philosophy reversed from

growing (+1.88% per month) to dropping (−6.37%). These results suggest a dissociation between political and foundational (i.e., psychological and philosophical) topics as the community engaged with psychological content individually and sought certainty by policing political norms in place of the structure Peterson had provided.

Impersonal policing of outgroups rather than reactive discharge. Reactive hostility did not spike toward specific outgroups. Overall Hostile Reactivity, which had been growing (+0.52% per month), dropped immediately (−2.03%), and Hostile Reactivity trended downward toward women (−1.17% per month) and in politics (reversing from +1.20% to −2.06% per month). Instead, Impersonal Policing grew toward both women (+3.67% per month) and LGBTQ+ individuals (+4.61% per month), and Group Baseline framing toward women reversed from declining (−1.41% per month) to growing (+3.98% per month). Women rose as a topic (+2.18%) with increased third-person referencing (+3.45%), and Third-Person Focus in political discussions reversed from declining (−0.65% per month) to growing (+0.98% per month). These findings suggest the community responded to internal collapse by converting outgroup hostility into impersonal norm enforcement rather than discharging it emotionally.

Hostility turned inward. As organized outgroup hostility declined, reactive hostility intensified within the community's own self-help discussions. Hostile Reactivity in failure-focused self-help, which had been declining before the event (−2.30% per month), spiked immediately (+7.08%) and continued as a trend (+2.33% per month), while dropping immediately toward positive self-help content (−3.38%, reversing from +0.66% per month). Hostile Reactivity in unclassified discourse followed the same pattern, reversing from decline (−2.33% per month) to spiking (+11.97%) and continuing to grow (+2.50% per month), while Anxious Self-Focus in unclassified discourse, which had been rising before the event (+1.43% per month), dropped immediately (−4.68%). Anxious Self-Focus in self-improvement discussions, which had been declining (−0.74% per month), began recovering (+0.83% per month). These findings suggest that without a stable epistemic anchor, the community's reactive energy turned inward, concentrating around failure and diffusing into unclassified venting rather than organizing into structured outgroup attack.

Collective outgroup pivot without cohesion. Overall Group Baseline rose immediately (+1.75%), and Group Baseline framing toward political outgroups reversed from declining (−3.85% per month) to spiking dramatically (+27.18%). Yet Unquestioning Affiliation showed no statistically significant change, and internal collective markers collapsed as Group Baseline dropped immediately in psychology (−8.10%) and failure-focused self-help (−4.33%, before recovering at +3.11% per month). These findings suggest the community formed strong collective opposition around political enemies, but without a stable internal anchor to defend, outgroup hostility discharged without converting into group cohesion.

III. Identity Reconstruction: Peterson's Reappearance

Peterson's public return in June 2020 challenged the community to reintegrate a destabilized leader and restore the internal anchor whose loss had lessened cohesion. The data indicate a mixed process as the community reorganized collective identity narrowly around political outgroups while collective framing collapsed across nearly every other domain. Unclassified discourse declined as a sustained trend (−1.13% per month), suggesting this reorganization proceeded by narrowing: conversation re-concentrated into the community's defined topics even as collective framing within most of them collapsed.

Continued avoidance of Peterson and internal topics. Peterson-related discussion dropped immediately (−4.19%), as did third-person referencing of him (−4.06%). Psychology

(−1.72%) and self-improvement (−1.77%) also declined immediately, with Group Baseline framing in psychology, which had been growing (+2.39% per month), collapsing (−11.55%) alongside Third-Person Focus in psychological discussions (−2.98%). These findings suggest that Peterson's return did not entirely restore his function as an epistemic anchor.

A large but transient pivot toward political outgroups. Political outgroup discussion surged immediately (+9.63%) alongside general politics (+1.07%), overall Third-Person Focus (+4.31%), Third-Person Focus in political outgroup discussions (+9.47%), and Group Baseline framing toward political outgroups (+35.05%), the largest single event-level effect in the dataset. Third-Person Focus also rose as a modest sustained trend within philosophical discussions (+0.68% per month). However, both Impersonal Policing (−6.09% per month) and Anxious Self-Focus (−0.87% per month) toward political outgroups declined as trends following the event. These findings suggest the community formed an intense, immediate collective orientation around political opposition that did not persist.

Broad collapse of collective framing. Although overall Group Baseline rose immediately (+2.77%, reversing a pre-event decline of −0.53% per month), collective framing simultaneously collapsed across internal topics: Group Baseline dropped in psychology (−11.55%), failure-focused self-help (−12.50%, reversing from +7.06% per month and continuing at −7.40% per month), discussions of women (−10.85%), and unclassified discourse (−9.81%). Impersonal Policing toward women also dropped immediately (−10.58%), while Anxious Self-Focus spiked toward women (+1.46%) and LGBTQ+ individuals (+1.09%) and dropped in self-improvement discussions (−5.48%). Hostile Reactivity in self-help reversed direction across subtopics (improvement: from −1.13% per month before the event to +1.37% per month after; positive: from +0.98% to −1.34% per month). Unquestioning Affiliation again showed no statistically significant change. These findings suggest that collective identity reorganized around political opposition alone while destabilizing everywhere else, marking the second consecutive event in which outgroup opposition failed to produce group cohesion.

IV. Intergroup Threat: Twitter Ban

In June 2022, Peterson was banned from Twitter following a transphobic tweet targeting Elliot Page, triggering substantial public backlash. Unlike during the intragroup threat events, Peterson was present and his ideological framework intact, so the community had both an enemy to oppose and a positive identity to defend. As such, this was the only threat event to produce significant growth in group cohesion.

Initial destabilization followed by gradual cohesion. The ban produced immediate destabilization across broad linguistic markers: Impersonal Policing dropped (−2.06%), Group Baseline dropped (−1.46%), Hostile Reactivity rose (+2.26%), and Third-Person Focus, which had been declining (−1.57% per month), spiked (+6.30%). Hostile Reactivity in unclassified discourse spiked sharply (+10.11%) before decaying (−2.73% per month), a transient flash of diffuse hostility that dissipated as the community's organized response took shape. Unquestioning Affiliation then built gradually as a sustained trend ($b = 0.045$ index units per month), the only significant Unquestioning Affiliation effect across the threat events.

Ideological retrenchment onto Peterson's core terrain. Alongside this destabilization, the community immediately consolidated around the ideological territory under attack. Impersonal Policing spiked in philosophy (+5.85%) and improvement-focused self-help (+7.38%), and Group Baseline framing in philosophy rose (+9.18%). Hostile Reactivity in improvement-focused self-help, which had been growing (+1.32% per month), collapsed immediately (−8.22%). Group Baseline framing in positive self-help reversed from declining

(−2.50% per month) to recovering (+2.49% per month), and improvement-focused self-help grew as a topic (+0.54% per month) alongside Third-Person Focus in improvement-focused (+0.73% per month) and positive self-help discussions (+0.83% per month). Group Baseline framing in psychology, which had been growing (+0.55% per month), instead collapsed immediately (−2.69%) and continued to decline (−0.73% per month). These findings suggest the community defended its self-help identity rather than lashing out through it.

Structured hostility. Despite the ban being triggered by a transphobic tweet, LGBTQ+ discussion declined as a trend (−0.66% per month), and third-person referencing in LGBTQ+ discussions, which had been growing (+0.52% per month), reversed and declined (−0.88% per month). Anxious Self-Focus toward political outgroups declined (−0.69% per month), and Hostile Reactivity in philosophy reversed from growing (+0.76% per month) to declining (−0.96% per month). These findings suggest a movement away from explicit outgroup labeling toward abstract, principle-based boundary enforcement, consistent with radicalization in which prejudice is legitimized through ideology rather than direct attack.

Sustained anxious and hostile political engagement. Anxious Self-Focus in political discussions reversed from declining (−0.68% per month) to spiking immediately (+1.72%) and continuing as a trend (+0.67% per month). Hostile Reactivity in politics similarly reversed from declining (−0.55% per month) to growing (+1.00% per month), and Hostile Reactivity in failure-focused self-help grew (+1.44% per month). These sustained trends suggest that the ban's political hostility became more organized over time as Unquestioning Affiliation built, consistent with the community refining rather than dissipating its opposition.

Discussion

This study provides a naturalistic, longitudinal test of the source model of identity threat (Greenaway & Cruwys, 2019) in a prominent online Manosphere community. Whereas prior work has examined intergroup and intragroup threats primarily in laboratory settings or cross-sectional surveys, eight years of naturalistic linguistic data from r/JordanPeterson show that the two threat types provoked distinct defensive responses. The intragroup threat was followed by gradual, impersonal norm enforcement toward women and LGBTQ+ individuals, reactive hostility that turned inward toward discussions of personal failure, and no detectable growth in group cohesion. The intergroup threat showed distinct responses, as explicit outgroup discussion declined while principle-based defense of the group's own identity content intensified, and group cohesion showed sustained growth. The findings suggest a refinement of the model: the source of a threat shapes the form a community's defense takes, i.e., what its hostility is aimed at, in what register, and whether that defense strengthens the group.

Our findings suggest that outgroup hostility converts into group cohesion only when the community simultaneously holds a stable internal anchor to defend. Intergroup threats activate this condition because the attack clarifies what the group stands for; intragroup threats suppress it because the source of the crisis is the anchor itself. This pattern extends the source model by specifying when outgroup hostility does and does not produce cohesion and is consistent with UIT's prediction that identity-clarifying structures reduce uncertainty. When the internal anchor was activated during the Twitter ban, Unquestioning Affiliation followed; when it was compromised, even the largest outgroup pivots left the community fragmented.

Threat Source Shapes the Form of Defensive Response

Following the 2018 release of *12 Rules for Life*, the community established a relatively benign foundation. Overt hostility toward women was suppressed, and discussions of feminized outgroups were marked by analytic distance rather than norm-enforcing or collective framing.

During this identity construction phase, members appeared to derive certainty and self-esteem from Peterson's philosophical framework itself, policing group boundaries through political and philosophical antagonism rather than through outgroup attack. Even at this stage, however, the seeds of externalization were visible when, in discussions of failure, members increasingly referenced other people rather than themselves, suggesting Peterson's frameworks were being turned outward to assign blame rather than inward for simple self-improvement.

When the community faced an intragroup threat, outgroup derogation emerged, but not in the hostile way that was predicted. Reactive hostility declined toward women and political outgroups and instead turned inward and outward into diffuse venting, intensifying around discussions of personal failure and within unclassified discourse. What grew toward outgroups was impersonal norm enforcement as Impersonal Policing rose steadily toward women and LGBTQ+ individuals. This pattern is consistent with compensatory scapegoating driven by self-uncertainty (Hogg, 2007); with the group's internal source of certainty compromised, members gradually constructed impersonal, rule-like derogation of the feminized outgroups already implicated in the community's crisis-of-masculinity narrative. Policing outgroup boundaries reasserted clear norms and categories of blame where reactive anger alone could not.

The intergroup threat produced a categorically different response, one in which no outgroup-directed linguistic marker significantly increased. Where the withdrawal produced a slow outward drift toward policing outgroups, the Twitter ban prompted immediate, structured opposition organized around the group's own identity content as explicit discussions of the LGBTQ+ outgroup declined even though the ban was triggered by a transphobic tweet, while directive, principle-based argument intensified on the ideological terrain under attack and political hostility organized gradually into sustained trends.

This shift from naming enemies to defending principles may be the more concerning trajectory; to the extent that exclusionary attitudes are reliant in abstract principles rather than reactive attacks, they may be more easily legitimized and more resistant to content moderation, a pattern that resembles descriptions of later-stage radicalization.

Outgroup Opposition Builds Cohesion Only Around an Internal Anchor

The central distinction in our findings concerns group cohesion. The two events that compromised the community's internal anchor, Peterson's withdrawal and his reappearance, produced the two largest collective outgroup pivots in the corpus (+27.18 and +35.05 percentage points in Group Baseline framing toward political outgroups), yet Unquestioning Affiliation showed no detectable change during either. The intergroup threat produced the only sustained growth in Unquestioning Affiliation across the threat events, carried out via increased use of Impersonal Policing and Group Baseline framing in philosophy and self-help, the identity content the ban had seemingly placed under attack.

This pattern extends the source model by specifying the conditions under which outgroup hostility produces group cohesion. The community needed something to defend, and the intergroup threat supplied it by clarifying what the group stood for. Intragroup threats suppressed this condition because the source of the crisis was the anchor itself. The finding also integrates naturally with uncertainty-identity theory (Hogg, 2007), as opposition without identity content offers no reduction in self-uncertainty, whereas defending a clear, positively valued framework offers both a problem and a solution.

Peterson's reappearance sharpens this account. If the internal anchor were simply the leader's presence, his return should have restored it. Instead, the community continued avoiding Peterson and its foundational topics, collective framing collapsed nearly everywhere, and

cohesion failed to rebuild. This suggests the group's philosophical anchor represents the credibility of the group's framework. As his withdrawal cast doubt on the group's promise of self-mastery, reappearance was insufficient for repair. Trust in an identity framework, once compromised from within, appears to require more than the removal of the original threat.

Revisiting the Predicted Responses

Returning to the expectations outlined in Table 3 clarifies what the source model did and did not anticipate. The predictions for the two non-threat events were partially observed. The book's release was followed by the predicted growth in cohesion and intense engagement with core identity topics, though boundary policing through political and philosophical antagonism emerged earlier than a purely benign reading would suggest. The reappearance prediction fared poorly: rather than the anticipated cohesion and repair, the community continued avoiding Peterson and its foundational topics, collective framing collapsed across internal domains, and cohesion did not rebuild. As we argue below, this failed prediction is theoretically informative.

The threat predictions were each somewhat accurate. Following the withdrawal, cohesion indeed failed to grow, and the community showed both of the anticipated alternative responses at once: outgroup derogation, in the form of impersonal policing of women and LGBTQ+ individuals, and intragroup conflict, in the form of reactive hostility concentrated in discussions of personal failure. Following the Twitter ban, the predicted growth in cohesion appeared, but the predicted derogation of outgroups, especially LGBTQ+ individuals, did not take its expected form, as explicit LGBTQ+ discussion declined, and opposition was instead routed through defense of the group's ideological terrain. In both cases, the source model correctly anticipated the cohesion outcome while underspecifying the form the defensive response would take. Specifying that form (inward-turned and impersonal under intragroup threat, but ideologically fortified under intergroup threat) is the refinement the present study contributes.

Scapgoating and the Reaffirmation of Masculinity

The macro-level trends situate these event dynamics within the community's broader drift. Over eight years, discourse moved steadily away from Peterson and psychology and toward politics and LGBTQ+ outgroup discussion, while reactive hostility toward women rose year over year and Unquestioning Affiliation climbed continuously. The community's trajectory thus mirrors the theorized radicalization pathway of the Manosphere, wherein a group formed to supply certainty and self-esteem to men experiencing identity precarity redirected its focus from self-improvement toward outgroup derogation that its narrative holds responsible. The event-level findings suggest that identity threats are inflection points that determine whether defensive responses took a raw, inward-turned form or a structured, impersonal one.

Future Directions

These findings offer theoretical and methodological directions for research on other figures and communities within the Manosphere, and for identity-based groups more broadly. Theoretically, research on neomasculinities and online extremism should attend to the source of identity threats rather than treating all threats as uniform accelerants. Whether a community fractures into hostility or becomes an entitative, norm-bound movement appears to depend on whether its internal identity content survives the threat. This distinction also carries practical implications: when cohesion fails to build and members' attachment is weakest, intragroup crises may offer times when counter-radicalization efforts and off-ramps are most likely to reach individual members, whereas intervention during periods of externally imposed threat may backfire by supplying the group with further evidence of persecution.

Methodologically, this study demonstrates the value of combining machine learning techniques to capture both the “how” and the “what” of group discourse: latent profile analysis classified the cognitive and affective style of posts, transformer-based topic modeling classified their content, and interrupted time series analysis tracked both against theoretically specified events. This pipeline captures linguistic patterns at a scale and granularity beyond human coding while remaining anchored to psychological theory, and it can be applied to any community with an archived textual record.

Limitations

Several limitations warrant caution. First, computational text analysis lacks the full context available to human readers, such as sarcasm and culturally specific slang. Although we implemented qualitative checks throughout coding and interpretation, some contextual subtleties may have been missed. Second, the exclusions required for reliable classification were substantial: of 209,500 collected posts, only 87,035 met the word count and classification criteria for analysis, and the retained sample likely overrepresents longer, more semantically coherent posts. Third, although event-specific findings were held to a three-criterion threshold to limit false positives, we did not apply a blanket correction for multiple comparisons, the threshold itself was adopted after initial inspection of the results rather than preregistered, and some individual findings may not replicate. Fourth, the design includes a single instance of each threat type, so threat source is confounded with other features distinguishing the time periods, including their historical moment, Peterson’s own ideological evolution, and the community’s changing size. The study is accordingly best read as a theoretically guided case study that generates a refined, testable account of threat responses rather than a definitive causal test. These findings also derive from a single group, and generalization to others should be tested.

Conclusion

This research supports an identity threat-based framework for understanding the Manosphere and the broader mechanics of online radicalization. The transition from seemingly benign self-improvement discourse into a politically hostile ecosystem was neither sudden nor uniform; rather, it unfolded through identity-relevant events whose psychological source shaped the language used in discussions of outgroups and whether those discussions strengthened the group. When young men experiencing identity precarity seek group membership for certainty and self-esteem, the resulting communities are vulnerable to threats from multiple sources, and the source of the threat shapes the form identity defense takes: the internal crisis turned hostility inward and toward impersonal policing of scapegoated outgroups, whereas the external attack redirected the community toward an ideologically fortified defense of its own identity. Only the threat that left the group a positive identity to defend converted opposition into cohesion. This conditional insight matters for intervention. Rather than treating polarization as a purely ideological problem, researchers and policymakers can target the identity needs that make extremist narratives appealing, and can time those efforts to the moments, such as internal crises, when a community’s hold on its members is weakest.

References

- Adelina, N., Chan, C. S., Takano, K., Yu, P. H. M., Wong, P. H. T., and Barry, T. J. (2023). The stories we tell influence the support we receive: Examining the reception of support seeking messages on Reddit. *Cyberpsychol. Behav. Soc. Netw.* 26, 823–834. doi: 10.1089/cyber.2023.0144
- Boyd, R. L., Ashokkumar, A., Seraj, S., and Pennebaker, J. W. (2022). *The Development and Psychometric Properties of LIWC-22*. Austin, TX: LIWC. <https://www.liwc.app>
- Branscombe, N. R., Ellemers, N., Spears, R., and Doosje, B. (1999). "The context and content of social identity threat," in *Social Identity: Context, Commitment, Content*, eds. N. Ellemers, R. Spears, and B. Doosje (Oxford: Blackwell Publishing), 35–58.
- Braun, V., and Clarke, V. (2006). Using thematic analysis in psychology. *Qual. Res. Psychol.* 3, 77–101. doi: 10.1191/1478088706qp063oa
- Brewer, M. B. (1991). The social self: On being the same and different at the same time. *Pers. Soc. Psychol. Bull.* 17, 475–482.
- Champagne, M. (2020). *Myth, Meaning, and Antifragile Individualism: On the Ideas of Jordan Peterson*. Exeter: Imprint Academic.
- Cohen, J. (1988). *Statistical Power Analysis for the Behavioral Sciences* (2nd ed.). Hillsdale, NJ: Lawrence Erlbaum Associates.
- Doosje, B., Ellemers, N., and Spears, R. (1995). Perceived intragroup variability as a function of group status and identification. *J. Exp. Soc. Psychol.* 31, 410–436.
- Eichstaedt, J. C., Smith, R. J., Merchant, R. M., Ungar, L. H., Crutchley, P., Preotiuc-Pietro, D., et al. (2018). Facebook language predicts depression in medical records. *Proc. Natl. Acad. Sci. U.S.A.* 115, 11203–11208. doi: 10.1073/pnas.1802331115
- Ellemers, N., Spears, R., and Doosje, B. (1997). Sticking together or falling apart: In-group identification as a psychological determinant of group commitment versus individual mobility. *J. Pers. Soc. Psychol.* 72, 617–626. doi: 10.1037/0022-3514.72.3.617
- Gee, J. P. (2005). *An Introduction to Discourse Analysis: Theory and Method* (2nd ed.). New York, NY: Routledge.
- Giles, M. W., and Evans, A. S. (1985). External threat, perceived threat, and group identity. *Soc. Sci. Q.* 66, 50–66.
- Gottzén, L., Fangen, K., and Qvotrup Jensen, S. (2026). Masculinity, masculinist politics, and political extremism. *J. Gend. Stud.* 1–7. doi: 10.1080/09589236.2026.2654169
- Greenaway, K. H., and Cruwys, T. (2019). The source model of group threat: Responding to internal and external threats. *Am. Psychol.* 74, 218–231. doi: 10.1037/amp0000321
- Grootendorst, M. (2022). BERTopic: Neural topic modeling with a class-based TF-IDF procedure. *arXiv [Preprint]*. doi: 10.48550/arXiv.2203.05794
- Habib, H., Srinivasan, P., and Nithyanand, R. (2022). Making a radical misogynist: How online social engagement with the manosphere influences traits of radicalization. *Proc. ACM Hum. Comput. Interact.* 6, 1–28. doi: 10.1145/3555551
- Han, Y., Pan, W., Li, J., Zhang, T., Zhang, Q., and Zhang, E. (2021). Developmental trend of subjective well-being of Weibo users during COVID-19: Online text analysis based on machine learning method. *Front. Psychol.* 12:779594. doi: 10.3389/fpsyg.2021.779594
- Hogg, M. A. (2007). "Uncertainty–identity theory," in *Advances in Experimental Social Psychology*, Vol. 39, ed. M. P. Zanna (San Diego, CA: Academic Press), 69–126. doi: 10.1016/S0065-2601(06)39002-8
- Hogg, M. A. (2014). From uncertainty to extremism: Social categorization and identity

- processes. *Curr. Dir. Psychol. Sci.* 23, 338–342. doi: 10.1177/0963721414540168
- Hogg, M. A., and Adelman, J. (2013). Uncertainty–identity theory: Extreme groups, radical behavior, and authoritarian leadership. *J. Soc. Issues* 69, 436–454. doi: 10.1111/josi.12023
- Institute for Strategic Dialogue (2022). *The “Manosphere”*. London: Institute for Strategic Dialogue. <https://www.isdglobal.org/explainers/the-manosphere/>
- Kimmel, M. (2017). *Angry White Men: American Masculinity at the End of an Era* (2nd ed.). New York, NY: PublicAffairs.
- Klein, O., Spears, R., and Reicher, S. (2007). Social identity performance: Extending the strategic side of SIDE. *Pers. Soc. Psychol. Rev.* 11, 28–45. doi: 10.1177/1088868306294588
- Koutsoumpis, A., Oostrom, J. K., Holtrop, D., Van Breda, W., Ghassemi, S., and De Vries, R. E. (2022). The kernel of truth in text-based personality assessment: A meta-analysis of the relations between the Big Five and the Linguistic Inquiry and Word Count (LIWC). *Psychol. Bull.* 148, 843–868. doi: 10.1037/bul0000381
- Lafrance, M., Menteth Wheelwright, I., and Ferguson, N. (2025). “The fuss is only just beginning”: Jordan Peterson and the mainstreaming of anti-gender ideology. *J. Gend. Stud.* 1–21. doi: 10.1080/09589236.2025.2478208
- Marques, J. M., Yzerbyt, V. Y., and Leyens, J.-P. (1988). The “black sheep effect”: Extremity of judgments towards ingroup members as a function of group identification. *Eur. J. Soc. Psychol.* 18, 1–16.
- Nesbitt-Larking, P. (2022). Constructing narratives of masculinity: Online followers of Jordan B. Peterson. *Psychol. Men Masculinities* 23, 309–320. doi: 10.1037/men0000394
- Nicholas, L., Clark, S., Agius, C., and Cook, K. (2026). Antifeminist, manosphere and right-wing extremist sentiment among men who use domestic and family violence: Masculinism, misinformation, and the justificatory logics of violence. *NORMA* 21, 96–115. doi: 10.1080/18902138.2024.2430513
- Panek, E. T. (2022). *Understanding Reddit* (1st ed.). New York, NY: Routledge. doi: 10.4324/9781003150800
- Pennebaker, J. W., Mehl, M. R., and Niederhoffer, K. G. (2003). Psychological aspects of natural language use: Our words, our selves. *Annu. Rev. Psychol.* 54, 547–577. doi: 10.1146/annurev.psych.54.101601.145041
- Peterson, J. B. (2022). *A Conservative Manifesto*. <https://www.jordanbpeterson.com/wp-content/uploads/2023/03/A-Conservative-Manifesto.pdf>
- Peterson, J. B., and The Daily Wire (2022). *DailyWire+ on Instagram: “Refusing to back down, @jordan.b.peterson calls out the ‘woke moralists’ as he deconstructs the tweet that put him in the censors’ crosshairs.”* <https://www.instagram.com/dailywireplus/reel/Cfe3MYggU7F/>
- Reimers, N., and Gurevych, I. (2019). “Sentence-BERT: Sentence embeddings using Siamese BERT-networks,” in *Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing and the 9th International Joint Conference on Natural Language Processing (EMNLP-IJCNLP)* (Hong Kong: Association for Computational Linguistics), 3980–3990. doi: 10.18653/v1/D19-1410
- Renström, E. A., and Bäck, H. (2024). Manfluencers and young men’s misogynistic attitudes: The role of perceived threats to men’s status. *Sex Roles* 90, 1787–1806. doi:

- 10.1007/s11199-024-01538-2
- Riek, B. M., Mania, E. W., and Gaertner, S. L. (2006). Intergroup threat and outgroup attitudes: A meta-analytic review. *Pers. Soc. Psychol. Rev.* 10, 336–353. doi: 10.1207/s15327957pspr1004_4
- Rios, K., Sosa, N., and Osborn, H. (2018). An experimental approach to intergroup threat theory: Manipulations, moderators, and consequences of realistic vs. symbolic threat. *Eur. Rev. Soc. Psychol.* 29, 212–255. doi: 10.1080/10463283.2018.1537049
- Rosenberg, J., Beymer, P., Anderson, D., Van Lissa, C. J., and Schmidt, J. (2018). tidyLPA: An R package to easily carry out latent profile analysis (LPA) using open-source or commercial software. *J. Open Source Softw.* 3:978. doi: 10.21105/joss.00978
- Schmid, K., and Muldoon, O. T. (2015). Perceived threat, social identification, and psychological well-being: The effects of political conflict exposure. *Polit. Psychol.* 36, 75–92. doi: 10.1111/pops.12073
- Sherman, D. K., Kinias, Z., Major, B., Kim, H. S., and Prenovost, M. (2007). The group as a resource: Reducing biased attributions for group success and failure via group affirmation. *Pers. Soc. Psychol. Bull.* 33, 1100–1112. doi: 10.1177/0146167207303027
- stuck_in_the_matrix, W. R. (2025). *Reddit comments/submissions 2005-06 to 2024-12* [Dataset]. Academic Torrents. <https://academictorrents.com/details/ba051999301b109eab37d16f027b3f49ade2de13>
- Tajfel, H., and Turner, J. C. (2000). "An integrative theory of intergroup conflict," in *Organizational Identity: A Reader*, eds. M. J. Hatch and M. Schultz (Oxford: Oxford University Press), 56–65. doi: 10.1093/oso/9780199269464.003.0005
- Tausczik, Y. R., and Pennebaker, J. W. (2010). The psychological meaning of words: LIWC and computerized text analysis methods. *J. Lang. Soc. Psychol.* 29, 24–54. doi: 10.1177/0261927X09351676
- Wang, W., Wei, F., Dong, L., Bao, H., Yang, N., and Zhou, M. (2020). MiniLM: Deep self attention distillation for task-agnostic compression of pre-trained transformers. *arXiv* [Preprint]. doi: 10.48550/arXiv.2002.10957
- Yan, Y., Li, J., Liu, X., Li, Q., and Yu, N. X. (2024). Identifying Reddit users at a high risk of suicide and their linguistic features during the COVID-19 pandemic: Growth-based trajectory model. *J. Med. Internet Res.* 26:e48907. doi: 10.2196/48907
- Zarate, D., Ball, M., Prokofieva, M., Kostakos, V., and Stavropoulos, V. (2023). Identifying self disclosed anxiety on Twitter: A natural language processing approach. *Psychiatry Res.* 330:115579. doi: 10.1016/j.psychres.2023.115579

Table 1. *Latent profiles with core LIWC variables and examples.*

Profile	N	LIWC Variables with z-scores	Example
Anxious Self-Focus	20838	High first-person singular (+1.43), low first-person plural (−0.28), low third-person (−0.26), high anxiety (+0.22), and low power (−0.32).	"I'm 28 years old, a Lyft driver, writing this out of a sense of desperation and sadness... I fear that I shall now be forced to choose the safe path and be condemned to a life of sacrificing my soul to the pursuit of stability over my own adventure."
Hostile Reactivity	2132	High anger (+4.88), low tone (−0.93), and low word count (−0.38).	"Elon Musk has dedicated his life to becoming the smartest, hardest working and capable person... The Redditors in general who are the most likely to circle jerk in those comments sections have absolutely nothing going on in their lives."
Third-Person Focus	7806	High third-person pronouns (+1.49).	"A lot of LGBTQ activists claim that they just want 'acceptance.' ... What I don't accept is that we need to constantly venerate them for being different and make public statements and iconography that effect. What I don't accept is that they need special privileges."
Group Baseline	52918	High use of first-person plural pronouns (+0.14).	"Do we really want Department of Justice not to prosecute a politician who deserves it just because someone suspects indictment might cause divisions in our society?"
Impersonal Policing	3341	Low word count (−0.41), negative tone (−0.5), low "I" focus (−0.89), high power (+3.35), low certitude (−0.28), low insight (−0.59), high moralization (+0.62).	"We need to defend against r/enoughpetersonspam because they brigade this sub, downvote everything, obfuscate and attack. Can somebody contact Reddit admins? Brigading is against Reddit rules."

Note. Word count and tone were not LPA indicators but are reported descriptively to aid profile interpretation.


Table 2. *Discursive themes with descriptions and examples.*

Theme	Description	Example
Jordan Peterson	Peterson as a public figure, including his appearances, life, and thoughts.	"More public speakers similar to Peterson? I'm looking for more talkers like Jordan Peterson to fill my Podcasts and youtube with..."
Self-help (negative)	Individuals' personal struggles and attempts to vent to others.	"The one thing I really wanted and liked to do is gone, due to my blundering... I try to remain hopeful but its very bleak... I've been reading your book " <i>12 rules for life</i> " and its really helpful."
Self-help (positive)	Individuals' personal triumphs, often to show acknowledge Peterson's role in their personal development.	"Hello Jordan... You helped me to find the confidence to stand on my own two feet and actively seek more responsibility which has led to propulsion in my life... I see you as one of, if not the most insightful humans I have ever read..."
Self-help (improvement)	Self-improvement methods, typically involving practical advice	"Help me make my life more challenging! I remeber that on one of the JRE epsidodes Peterson mentioned... that we won't have any challenges to complete and feel fullfilled by..."
Politics	Politically charged events or specific political beliefs.	"...demoralize an entire generation of American students using marxist-indoctrinated professors. Proof that the our current culture of people clinging to ideological dogma rather than using fact and reason is a literal marxist plot."
Psychology	Psychology as a field of general interest, for self-help purposes or knowledge.	"...are there certain archetypes that are in a sense meta-archetypes? And by that I mean: certain archetypes maybe are not expressed in an early stage of human development but emerge later..."
Philosophy	Philosophy as a field of general interest, for self-help purposes or knowledge.	"What is the argument for western ethics having a basis in the Bible?... It's probably very difficult to find a definitive answer as to where western ethics come from..."
Outgroups (women)	Women and/or feminism, including discussions of men or masculinity explicitly defined in opposition	"I was sitting with a few female acquaintances... One thing that bothered me is how selective/picky they were in terms of the guy they choose... I've remembered what Dr. Peterson said about how selective women are... we(men) fear rejection, how you(women) set the standards too high."
Outgroups (LGBTQ+)	LGBTQ+ people and/or movements in support or opposition to LGBTQ+ rights.	"...lgbtq community who were trying to force the idea that it is okay to give your children hormone blockers...corrupting the minds of innocent children with sickening ideas..."
Outgroups (political)	Other outgroups, including racial, ideological, or economic outgroups.	"...PC culture forced on us Westerners. It is the very destruction of what constitutes us, what makes us strong. I feel we are under siege... PR agency trolls spewing enforced PC viewpoints..."

Table 3. *Identity-relevant events and their expected responses.*

Event	Date	Event Type	Expected Response
<i>12 Rules for Life</i> release	January 2018	Positive	High group cohesion, discussions of core group identity topics
Peterson's withdrawal	November 2019	Intragroup threat	Low group cohesion, either outgroup derogation or intragroup conflict
Peterson's reappearance	June 2020	Positive	High group cohesion, discussions of core group identity topics, repair after previous conflict
Twitter ban and public criticism	June 2022	Intergroup threat	High group cohesion, derogation of outgroups (especially LGBTQ+)

Figure 1
Discourse Profile Composition by Post Theme.



Required Statements

Acknowledgements

The authors thank Christopher Forth for theoretical guidance and feedback on early drafts. Additional gratitude is extended to the University of Kansas's Center for Undergraduate Research & Fellowships for funding support, as well as the Kansas Data Science Consortium for assistance with coding. We also acknowledge the use of Anthropic's Claude Opus 4 models for language polishing.

Funding

This work was supported by an Undergraduate Research Award from the University of Kansas. The University of Kansas also provided funding, via institutional partnership, for open access publication fees. The funders had no role in the study design, data collection and analysis, or preparation of the manuscript.